



AI Application Developer - Technical Assessment

Overview:

- **Expected Duration:** Around 3-5 hours
- **Tools allowed:** ChatGPT, Claude, Gemini, Cursor, Replit, etc. (any AI coding tool and/or LLM)
- **Outcomes:**
 - Test applicant's proficiency with AI coding tools
 - Gauge applicant's approach to open-ended business problems
 - Evaluate applicant's AI systems design thinking
- **Deliverables:**
 - Link to a working POC and code
 - Write-up (2 pages max) explaining approach, results, and next steps

Scenario:

You've been asked by the **Head of Ecommerce** to explore whether an AI-powered experience could help customers find the right product more quickly and confidently on edible.com.

Edible Arrangements has a large catalog of products, and customers often:

- Don't know where to start
- Struggle to compare similar items
- Abandon sessions without purchasing

Leadership wants to test whether AI can **support better purchase decisions** — without misleading customers or overstepping trust.

Your Task:

Design and build an **AI-assisted product discovery experience** that helps users move toward a confident purchase decision.

This exercise is intentionally open-ended. There is **no single correct solution**.



Core Functionality:

- An LLM-powered system that accepts user intent and helps users move toward a confident purchase decision (e.g., recommendations, comparisons, clarifying questions, etc.)
 - Examples: Chatbot, AI-powered product search, etc. (the decision you make here is part of the test, bonus points for novel ideas)
- Integration with Edible catalogue via <https://www.ediblearrangements.com/api/search/> API
 - Takes POST requests with a `keyword` parameter
 - Returns top 50 SKUs (plus metadata) that match keyword
 - Sample Code: [Link](#)

These are the basics – you are encouraged to add/expand on functionality as you see fit. Keep in mind this is a proof-of-concept, so you do not need to overly optimize for speed or cost.

Key Questions to Explore

- Who is the primary user?
- When should the AI stop guiding and let the user decide?
- How do you prevent hallucinated claims?
- How do you handle uncertainty or ambiguous preferences?
- Prompt structure or modularity
- Model selection and cost/latency tradeoffs

Evaluation Criteria (100 points)

- **Core Functionality** (20): Does the solution work, and is the scope appropriate?
- **Use of AI Tools** (15): How effectively did you use AI tools to accelerate development?
- **AI System Design** (15): Is the GenAI system well-structured and reliable?
- **Output Quality & UX** (15): Are LLM outputs relevant, helpful, and appropriately toned?
- **Business/Product Understanding** (15): Does the solution align with the business goals?
- **Taste & Judgment** (10): Did you make good assumptions and thoughtful trade-offs?
- **Next Steps** (10): What would you improve before shipping this to real customers?



Submission: Reply to this email with the deliverables within the specified time limit.

Next Steps: If your assessment passes the initial screening, you will be invited to a 30-minute video call to walk through your solution and participate in a technical interview on AI engineering fundamentals.